

THE CANDOR PROJECT

Methodology and Evidentiary Standards

A companion document to "Epistemic Accountability and the Constitution: A Framework for Professional Truth Standards in Public Life"

Erik Muir

Founder, The Candor Project
erik@thecandorproject.org

Version 1.0 | June 2026

The author used AI tools (Claude, Anthropic) for research assistance and critical feedback during the drafting of this paper.

I. Purpose and Scope

This document is the operational companion to the white paper, *Epistemic Accountability and the Constitution: A Framework for Professional Truth Standards in Public Life*. Where the white paper establishes the legal architecture and constitutional theory of the Candor framework, this document operationalizes it by defining with precision the standards, procedures, and evidentiary requirements that govern every finding the oversight body produces. The two documents are designed to be read together; the constitutional defensibility of the framework depends in part on the rigor of the methodology described here.

The methodology document serves three distinct functions. First, it defines the boundaries of the oversight body's jurisdiction: the precise line between claims the body may review and the vast domain of political speech, opinion, and contested judgment it may not. Second, it specifies the evidentiary standards that govern how the body reaches its findings, ensuring that every conclusion is grounded in a documented, reproducible process rather than institutional discretion. Third, it establishes the procedural architecture through which findings are made, challenged, and resolved, including the appeals process that functions as the primary structural check on the body's own authority.

The oversight body's mandate is narrow by design. It applies only to empirically falsifiable factual claims made by covered officials and covered media entities in their official or broadcast capacities. It does not extend to opinion, political judgment, policy advocacy, forward-looking predictions, or claims in domains where the expert community remains genuinely divided. These exclusions are not concessions to political convenience; they are the principled boundaries that make the framework constitutionally defensible and institutionally credible. A methodology that reaches contested political territory is not a neutral scorekeeper but a political actor, and these boundaries exist to ensure the body remains the former.

The framework operates on two tracks, each with a distinct legal basis and enforcement architecture. The politician track covers the President, Vice President, Cabinet secretaries, members of Congress, federal judges, and

Supreme Court justices, applying to statements made in their official capacities. The media track covers public broadcasters receiving federal funding and licensed broadcasters operating on the public electromagnetic spectrum, applying to statements made in their broadcast capacities. The threshold evidentiary standards — the double filter described in Section II — apply uniformly across both tracks. The knowledge requirements and enforcement architecture differ between tracks and are addressed separately in Sections IV and V.

This document does not address the oversight body's appointment structure, funding mechanism, or governance, all of which are detailed in the white paper. It addresses only the intellectual and procedural core of the body's function: how it decides what is false, how it documents that determination, and how that determination can be challenged.

II. The Double Filter

Every claim the oversight body considers must pass through a two-part test before it becomes eligible for review. Both conditions must be satisfied simultaneously; a claim that meets one but not the other falls outside the body's jurisdiction. This threshold requirement is the methodology's primary defense against overreach, and its application is mandatory regardless of the political salience of the claim or the public profile of the official involved.

The first condition is empirical falsifiability. A claim is empirically falsifiable only if it asserts a specific, verifiable fact that can be confirmed or refuted by objective evidence. This means the claim must have a determinable truth value that does not depend on the assessor's values, political commitments, or interpretive framework. A statement that an unemployment rate reached a specific percentage on a specific date is empirically falsifiable; a statement that an economic policy was good or bad for working Americans is not. A statement that a specific document was produced on a specific date is empirically falsifiable; a statement that an administration was corrupt is not. The falsifiability condition functions as a filter against the most common category of political speech: evaluative claims dressed in factual language.

The second condition is an unambiguous evidentiary record. Even a claim that is empirically falsifiable does not qualify for review unless it is contradicted by evidence that is not subject to serious dispute among relevant experts. This condition does two things simultaneously. It ensures the body operates only where the factual ground is firm, and it prevents the body from being drawn into empirical debates where reasonable experts genuinely disagree. A claim contradicted by peer-reviewed scientific consensus, official government records, or the speaker's own prior documented statements meets this condition. A claim that is probably false, or false according to some credible analysts but not others, does not. The unambiguous threshold is deliberately high; the framework accepts that some false statements will fall below it rather than risk the body adjudicating genuinely contested empirical territory.

Together, these two conditions form the double filter. A claim must be both falsifiable in principle and contradicted by an unambiguous record in practice. The filter is applied at the threshold of review, before any evidentiary analysis begins, and a claim that fails either condition is returned without a finding. This is not a procedural technicality; it is the structural mechanism that keeps the body's mandate narrow and its findings defensible.

Three categories of statements warrant specific attention at the threshold stage because they frequently resemble falsifiable factual claims without meeting the standard.

Comparative superlatives present as factual but are rarely falsifiable. A statement that a given period represented the greatest economy in American history, or the most corrupt administration, asserts a ranking that depends entirely on which metrics are selected and how they are weighted. Because reasonable analysts applying different but legitimate frameworks will reach different conclusions, these statements do not have a determinable truth value and fail the falsifiability condition regardless of how confidently they are asserted.

Statements of intent and attribution present a related problem. A claim that a political opponent intends to destroy a program, or that a policy is designed to benefit a specific class of donors, asserts a fact about motivation rather than about the observable world. Intent is not directly observable and cannot be confirmed or refuted by the categories of evidence the methodology recognizes. These statements fail the falsifiability condition even when the attributed intent seems obvious or well-supported by circumstantial evidence.

Statistical claims with genuinely contested methodologies occupy a third threshold category. Where the underlying data is not in dispute but reasonable experts disagree about the appropriate analytical framework, the evidentiary record is not unambiguous even if one interpretation is more widely held. The body does not adjudicate methodological disputes among experts; it applies findings only where the dispute has been resolved to the satisfaction of the relevant expert community.

III. Exclusive Categories

The double filter establishes the threshold conditions for review. This section defines the categories of statements that fall outside those conditions, providing the precision necessary for consistent application. Each exclusion is stated as a general principle, followed by the analytical reasoning that supports it, and illustrated with examples drawn from both tracks and both sides of the partisan divide. These exclusions are exhaustive in category but not in example; the oversight body is expected to apply the underlying reasoning to novel cases rather than pattern-match against the illustrations provided here.

Opinion and Value Judgment

Statements that express a preference, assign value, or characterize a state of affairs in terms that depend on the speaker's normative framework are excluded from review regardless of how confidently or specifically they are asserted. The exclusion applies even when the underlying facts are not in dispute, because the evaluative conclusion drawn from those facts is not itself a factual claim. A politician who asserts that a tax policy is unfair, that an administration's foreign policy was a failure, or that a piece of legislation represents the greatest threat to American prosperity in a generation is making a claim about values and priorities, not about the observable world. The oversight body has no evidentiary basis for adjudicating such claims and no legitimate authority to do so.

This exclusion requires care at its edges. An evaluative conclusion can be bundled with a specific factual premise in ways that make the whole statement appear falsifiable. When that occurs, the body may review the specific factual premise if it independently meets the double filter, but the evaluative conclusion attached to

it remains outside the body's jurisdiction. The finding, if any, attaches to the factual component only.

Forward-Looking Predictions

Statements about future states of the world are excluded from review because they cannot be contradicted by an existing evidentiary record at the time they are made. A prediction that a proposed policy will create millions of jobs, that a trade agreement will reduce the deficit, or that tax cuts will generate sufficient economic growth to pay for themselves is a claim about a future that has not yet occurred. The subsequent failure of a prediction to materialize does not retroactively satisfy the unambiguous evidentiary record condition, because the standard applies at the time of the statement, not in hindsight.

This exclusion applies even when the prediction contradicts the preponderance of expert forecasting available at the time. The methodology does not treat expert consensus about the future as an unambiguous evidentiary record of a present fact; forecasting involves genuine uncertainty that the framework respects. A politician who asserts that tax cuts will pay for themselves despite contrary CBO projections is making a contestable economic prediction, not a falsifiable factual claim, and the existence of contrary projections does not transform the prediction into a reviewable statement. Similarly, a politician who characterizes an inflation trend as transitory is offering an economic forecast rather than asserting a verifiable present fact, and the subsequent persistence of inflation does not constitute a finding of falsity under this methodology.

Contested Scientific and Empirical Territory

Statements in domains where the relevant expert community remains genuinely divided are excluded from review because the unambiguous evidentiary record condition cannot be satisfied where serious dispute exists among experts. This exclusion is not a concession to manufactured controversy; it applies only where the disagreement is substantive and the participants are operating in good faith within established disciplinary frameworks. The oversight body is not equipped to resolve active scientific debates, and any attempt to do so would compromise both its mandate and its credibility.

The boundary of this exclusion requires precise application. The existence of minority dissent within an expert community does not render an evidentiary

record ambiguous if the mainstream consensus is overwhelming and the dissenting position lacks support in peer-reviewed literature. Conversely, a claim that attributes a specific weather event definitively to climate change falls within this exclusion even though the broader reality of anthropogenic climate change does not, because the science of attribution — linking specific events to long-term trends — involves probabilistic reasoning that the expert community has not resolved to the level of certainty the unambiguous threshold requires.

Rhetorical Hyperbole

Statements that deploy exaggeration for emphasis rather than to assert a literal fact are excluded from review on the grounds that they do not meet the falsifiability condition. A politician who describes an opponent's record as a disaster, characterizes a legislative session as the most productive in history, or claims that a proposed policy will destroy the country is using language whose communicative function is evaluative and emphatic rather than factually assertive. These statements are understood by speakers and audiences alike as rhetorical rather than literal, and treating them as falsifiable factual claims would require the body to apply a standard of literal construction that does not reflect how political speech actually functions.

The distinction between rhetorical hyperbole and a falsifiable factual claim is not always self-evident, and the body must apply a contextual standard: would a reasonable audience understand the statement as a literal factual assertion or as emphatic characterization? A statement that a specific program costs twice what the government claims is a literal factual assertion subject to review. A statement that the same program is a complete waste of taxpayer money is emphatic characterization that is not.

Statements of Intent and Attributed Motivation

Statements that characterize the intent, motivation, or purpose behind another actor's conduct are excluded from review because intent is not directly observable and cannot be confirmed or refuted by the evidentiary categories the methodology recognizes. A claim that an opposing party intends to eliminate a government program, that a legislative provision was designed to benefit a specific donor class, or that an administration's policy choices reflect a deliberate strategy to harm a particular community asserts a fact about internal mental states rather than about the observable world. Even where circumstantial

evidence makes an attribution of intent plausible or widely shared, the evidentiary record does not meet the unambiguous threshold required for a finding.

This exclusion applies symmetrically: it protects statements attributing malicious intent and statements attributing benign intent alike. The body does not adjudicate questions of political motivation regardless of which direction the attribution runs.

Scope Limitations

Two additional exclusions operate not on evidentiary grounds but on jurisdictional ones, and are worth stating explicitly to prevent category confusion.

Campaign rhetoric and private communications fall outside the politician track's jurisdiction regardless of their content, because the track applies only to statements made in an official capacity. A candidate's claims during a campaign, a politician's private correspondence, and personal social media activity that does not deploy the weight of official office are not subject to review even if they would otherwise meet the double filter. The capacity determination precedes the evidentiary analysis; a statement that fails the capacity threshold is not returned without a finding but is returned without jurisdiction.

Print and digital media fall outside the media track's jurisdiction regardless of their content, because the track applies only to public broadcasters receiving federal funding and licensed broadcasters operating on the public electromagnetic spectrum. An outlet such as a national magazine or an online publication may publish statements that would meet the double filter if made by a covered entity, but the absence of a licensing or funding relationship with the federal government places those statements beyond the framework's reach. This is not an endorsement of the accuracy of such outlets; it is a recognition that the legal basis for accountability differs categorically from the basis that governs licensed and public broadcasters.

IV. Evidentiary Standards

IV-A: Politician Track

A claim that passes the double filter and falls outside the exclusion categories established in Section III becomes eligible for evidentiary review. Eligibility is not a finding; it is the threshold at which the body's substantive analysis begins. The evidentiary review determines whether the factual record supporting a finding meets the burden of proof the methodology requires.

The framework adopts a burden of proof approximating clear and convincing evidence, a standard drawn from established civil law that requires substantially more than a preponderance of the evidence but does not reach the criminal threshold of beyond a reasonable doubt. This standard was selected deliberately. A preponderance standard — more likely than not — is insufficient for a mechanism whose findings feed into professional disqualification proceedings; the consequences are serious enough to demand a higher bar. A beyond reasonable doubt standard, conversely, would render the body nearly inoperative, since official false statements are rarely accompanied by direct admissions. Clear and convincing evidence occupies the appropriate middle ground: demanding enough to prevent weaponization, achievable enough to function.

To satisfy this burden, a finding must be grounded in at least one of four authoritative evidentiary categories. These categories are not a menu from which the body selects the most convenient option; each represents a class of evidence whose reliability and independence from the body's own judgment has been established through existing institutional practice.

The first category is peer-reviewed scientific or empirical research. This encompasses findings published in peer-reviewed journals and supported by the mainstream consensus of the relevant expert community. A single study, even a peer-reviewed one, does not satisfy this category unless its findings have been replicated or are consistent with the broader body of research. The body does not conduct its own scientific analysis; it assesses whether the peer-reviewed record is sufficiently settled to support a finding.

The second category is official government data. This encompasses statistics, records, and findings produced by federal agencies, congressional offices, and official international bodies whose methodologies are publicly documented and

subject to independent audit. Data produced by partisan advocacy organizations, think tanks with disclosed political affiliations, or governmental bodies whose independence is subject to serious dispute does not satisfy this category regardless of the accuracy of the underlying numbers.

The third category is the speaker's own prior documented statements. Where an official has made a specific factual assertion in a prior official capacity statement that directly contradicts a current assertion, the prior statement constitutes evidence against the current one. This category applies only to statements made in official capacity; prior campaign statements, private communications, and personal social media activity do not qualify as evidentiary inputs under this category for the same jurisdictional reasons they fall outside the politician track's coverage.

The fourth category is an established historical record verified by multiple independent sources. This encompasses facts whose accuracy has been confirmed through independent documentary evidence, contemporaneous records, or the convergent findings of multiple credible investigative bodies. A fact verified by a single source, however credible, does not satisfy this category; the independence and multiplicity of sources is the condition that distinguishes an established historical record from a contested account.

The methodology explicitly identifies several classes of evidence that do not satisfy the burden of proof regardless of their surface plausibility. Non-replicated studies, even when peer-reviewed, do not meet the first category's replication requirement. Contested interpretations of ambiguous data do not meet the unambiguous threshold established by the double filter. Partisan advocacy findings, even when produced by credible organizations, do not meet the independence requirement of the second or fourth categories. And evidence that is probative but not conclusive — the kind that establishes that a statement was probably false rather than clearly and convincingly false — does not satisfy the burden even when the preponderance of the record points in one direction.

Every finding issued by the oversight body must identify the specific evidentiary category or categories on which it relies, cite the specific sources within those categories, and explain why the evidence meets the clear and convincing standard. This requirement is not administrative formality; it is the mechanism that makes every finding specifically appealable and subjects the body's reasoning to continuous public scrutiny.

IV-B: Media Track

The evidentiary standards applied to the media track are structurally identical to those of the politician track: the same four authoritative categories, the same clear and convincing burden, and the same documentation requirements for issued findings. The differences between the two tracks at this stage are not differences of standard but of application, reflecting the distinct nature of broadcast entities and the correction record infrastructure the framework requires them to maintain.

The most significant practical distinction concerns the role of the correction record. The framework requires covered broadcasters to maintain standardized, publicly accessible correction policies as a condition of their licensing or funding relationship. Where such a record exists, it becomes an evidentiary input in its own right. A broadcaster that has issued a formal correction of a specific claim has produced a documented acknowledgment that the original statement was false; that acknowledgment is treated as evidence equivalent to the speaker's own prior documented statements under the third evidentiary category. Conversely, a broadcaster that has failed to issue a correction for a claim later shown to be false, despite possessing information that should have prompted one, creates an evidentiary trail relevant to the knowledge analysis addressed in Section V-B.

The second distinction concerns the unit of analysis. The politician track applies to specific statements made by identified officials in their official capacities. The media track applies to statements made in broadcast capacity, which may involve multiple contributors — anchors, correspondents, producers, and editorial decision-makers — whose individual roles in producing a false statement differ. For the purposes of evidentiary review, the relevant unit is the broadcast entity rather than the individual contributor, and the evidentiary analysis addresses the outlet's institutional knowledge and editorial process rather than the mental state of any single employee. This distinction has consequences for the knowledge analysis and is addressed further in Section V-B.

The four evidentiary categories apply to media track findings with one additional consideration. Broadcasters frequently rely on sourced reporting, in which a statement's accuracy depends on the reliability of a confidential or named source rather than on the broadcaster's independent verification. Sourced-but-wrong reporting — where a broadcaster accurately reported what a source told them, but the source's information proved false — does not satisfy the clear and convincing

standard for a finding against the broadcaster, because the evidentiary record does not establish that the broadcaster's statement was false at the time it was made in the relevant sense. The body's analysis in sourced reporting cases focuses on whether the broadcaster's verification process met the standards its own publicly documented editorial policies require, not on whether the underlying information later proved inaccurate.

V. Knowledge Requirements

A finding requires not only that a statement be contradicted by an unambiguous evidentiary record, but that the covered official or entity had a sufficient relationship to that record at the time the statement was made. This requirement exists because the framework is designed to identify and document deliberate or negligent misrepresentation, not honest error. A public official who states something false while genuinely lacking access to contradicting information has failed a different standard than one who states something false while possessing documentation that directly contradicts it. The knowledge requirement is the mechanism that distinguishes these cases.

The framework adopts a two-part knowledge standard: actual knowledge and constructive knowledge. Either is sufficient to support a finding; neither requires direct evidence of intent to deceive.

V-A: Politician Track

Actual knowledge is established when the evidentiary record demonstrates that the official possessed specific information contradicting their statement at the time it was made. This does not require a direct admission; it can be established through documentary evidence such as briefing materials, agency reports, classified assessments available to the official in their capacity, or prior statements by the official themselves that reflect awareness of the contradicting information. The standard is not that the official read and absorbed every document available to them, but that the contradicting information was within their documented informational environment in a form that a person in their position would reasonably be expected to have encountered.

Constructive knowledge is established when the contradicting information was sufficiently available and prominent that a reasonable person exercising the duties of the official's specific role would have been aware of it. This standard is calibrated to the official's position rather than applied uniformly. The President and Cabinet secretaries are presumed to have constructive knowledge of information routinely surfaced through the daily intelligence briefing process, agency reporting, and the public scientific record in domains directly relevant to their portfolios. Members of Congress are presumed to have constructive knowledge of information in the public record, official government data, and testimony entered into the congressional record in committees on which they serve. Federal judges are presumed to have constructive knowledge of the evidentiary record as established in proceedings before them.

The constructive knowledge standard does not require that an official have sought out contradicting information; it requires only that the information was available through the normal exercise of their official duties. An official who makes a false statement about economic data that contradicts published Bureau of Labor Statistics figures cannot defeat a finding by demonstrating they did not personally consult those figures, because consulting publicly available official data is a reasonable expectation of the role. Conversely, an official who makes a false statement about a matter outside their portfolio and outside the normal scope of their informational environment may defeat a finding on constructive knowledge grounds even if the contradicting information existed in the public record.

One category of statement warrants specific treatment under the knowledge standard: statements that contradict the official's own prior documented statements made in official capacity. Where such a contradiction exists, actual knowledge is presumed rather than inferred, because a person cannot lack awareness of their own prior assertions. This presumption is rebuttable only by evidence that the official's position changed on the basis of new information that postdates the prior statement, in which case the body's analysis addresses whether the new position is itself supported by the evidentiary record.

V-B: Media Track

The knowledge standard applied to the media track is structurally parallel to the politician track but adapted to reflect the institutional rather than individual unit of analysis established in Section IV-B. Because findings on the media track

attach to the broadcast entity rather than to individual contributors, the knowledge analysis addresses the outlet's institutional knowledge and editorial process rather than the awareness of any single employee.

Actual institutional knowledge is established when the evidentiary record demonstrates that the outlet's editorial process had access to information contradicting the broadcast statement at the time it was made. This encompasses information available to producers, editors, fact-checkers, or senior editorial decision-makers in the normal course of preparing the relevant broadcast. The Fox News and Dominion Voting Systems case illustrates this category: internal communications produced during litigation documented that hosts, producers, and executives possessed specific information contradicting on-air claims about voting machine fraud at the time those claims were being broadcast. The evidentiary record in that case established actual institutional knowledge through documentary evidence of a kind that rarely exists in such completeness, which is why it functions as the clearest available example of a qualifying media track finding.

Constructive institutional knowledge is established when the contradicting information was sufficiently available and prominent that a broadcast outlet exercising reasonable editorial diligence would have encountered it. The standard here is the outlet's own publicly documented editorial policies. A broadcaster whose stated verification standards require independent corroboration of specific factual claims before broadcast is held to those standards; a finding of constructive knowledge is supported when a claim was broadcast without the verification the outlet's own policies require and the contradicting information was available through the verification process those policies mandate. This is why the framework's requirement that covered broadcasters maintain and publish standardized editorial policies is not merely a transparency condition — it creates the evidentiary baseline against which constructive knowledge is assessed.

The CBS News and Rathergate case illustrates the constructive knowledge analysis at its most consequential boundary. CBS broadcast a report whose central factual claim — the authenticity of specific documents — was contradicted by forensic evidence that typography experts raised almost immediately after broadcast. The constructive knowledge question in that case is whether CBS's pre-broadcast verification process should have surfaced the forensic vulnerabilities in the documents, given that the anomalies were identifiable through analysis available to experts in the relevant field. The post-broadcast

period, during which CBS initially resisted the forensic evidence before ultimately retracting, is relevant to the knowledge analysis because it documents the point at which constructive knowledge became actual knowledge and the outlet's editorial response to that transition. This case does not present the clean evidentiary record of the Dominion litigation, which is why it belongs in the constructive knowledge analysis rather than the actual knowledge category, and why it is appropriately treated as a harder case than the Fox example despite both being qualifying findings.

The sourced reporting limitation established in Section IV-B interacts with the knowledge standard in one specific way that warrants explicit treatment. Where a broadcaster accurately reported a source's claim and the source's information proved false, the knowledge analysis focuses on whether the broadcaster's verification process met its own documented editorial standards. If those standards were followed and the source's falsity was not reasonably detectable through the verification process the standards require, neither actual nor constructive knowledge is established and therefore no finding would be issued. If the standards were not followed, or if information contradicting the source was available through the verification process the outlet's own policies mandate and was not pursued, constructive knowledge may be established on the basis of the editorial process failure rather than on the basis of the source's underlying falsity.

VI. Qualifying Examples

The examples in this section illustrate the methodology's application to statements that meet the double filter, fall outside the exclusion categories, satisfy the evidentiary standard, and establish the required knowledge. They are drawn from both tracks and both sides of the partisan divide, and are selected because the evidentiary record in each case is sufficiently documented to support analysis. They are not presented as formal findings of the oversight body; they are analytical illustrations of how the methodology functions when applied to real cases.

Each example follows a consistent structure: the statement, the capacity in which it was made, the evidentiary record that contradicts it, and the knowledge analysis that completes the finding. This structure mirrors the documentation format the oversight body would use for an actual finding, and is presented in

that form deliberately so that the methodology's application is visible at every step.

Politician Track

Example 1: "If you like your health care plan, you can keep it"

President Obama made this assertion repeatedly in official capacity statements — including addresses to Congress and formal press appearances — during the legislative debate over the Affordable Care Act and after its passage. The statement is a specific, falsifiable factual claim about the consequences of the legislation for existing insurance arrangements, not a prediction about future policy outcomes or an evaluative characterization of the law's merits.

The evidentiary record contradicts the statement on the basis of official government data and the legislative record. Internal administration analyses, subsequently documented through congressional oversight proceedings, indicated that a significant proportion of existing individual market plans would not meet the ACA's minimum coverage requirements and would therefore be discontinued. The administration possessed this analysis during the period in which the statement was made in official capacity.

The knowledge analysis establishes actual knowledge. Documentary evidence produced through congressional oversight demonstrates that administration officials had access to actuarial projections contradicting the categorical assertion at the time it was being made publicly. The statement was not a forward-looking prediction about uncertain policy consequences; it was a present-tense categorical assertion made while the administration possessed information indicating the assertion was not accurate for a material portion of the affected population.

Example 2: Hydroxychloroquine "100% effective"

President Trump stated at an official White House briefing that hydroxychloroquine was "100% effective" against COVID-19. The statement is a specific, falsifiable factual claim about the efficacy of a medical treatment, made in official capacity at a formal government briefing.

The evidentiary record contradicts the statement on the basis of peer-reviewed scientific research and contemporaneous FDA guidance. At the time of the

statement, the FDA had issued guidance explicitly noting that the evidence for hydroxychloroquine's efficacy against COVID-19 was insufficient to support the claim being made, and no peer-reviewed research supported a finding of complete efficacy. The claim was contradicted not by subsequent research but by the scientific record as it existed at the time of the statement.

The knowledge analysis establishes actual knowledge. The White House medical advisory structure, including the Coronavirus Task Force, had access to FDA guidance and the existing research literature. The statement was made at a briefing at which medical advisors were present whose documented positions contradicted the specific efficacy claim being asserted.

Example 3: Inauguration crowd size

White House Press Secretary Sean Spicer stated at an official White House press briefing that President Trump's inauguration drew the largest audience ever to witness an inauguration, both in person and around the globe. The statement is a specific, falsifiable factual claim about a measurable quantity, made in official capacity at a formal government briefing.

The evidentiary record contradicts the statement on the basis of official government data and an established record verified by multiple independent sources. Washington DC Metro ridership data for inauguration day, published by the Washington Metropolitan Area Transit Authority, showed substantially lower ridership than comparable figures for prior inaugurations. Aerial and ground-level photographic documentation, independently verified, showed crowd density inconsistent with the claim. National Park Service assessments were consistent with these figures.

The knowledge analysis establishes constructive knowledge. The Metro ridership data was publicly available at the time of the briefing. A White House Press Secretary exercising the normal informational duties of that role would reasonably be expected to have encountered publicly available quantitative data directly relevant to a specific factual claim being made at an official government briefing.

Example 4: Student loan forgiveness "passed by Congress"

President Biden stated on multiple occasions in official capacity settings, including formal press appearances and public addresses, that his student loan

forgiveness program had been passed by Congress. The statement is a specific, falsifiable factual claim about the legislative origin of a federal program, not a characterization of the program's merits or a prediction about its legal durability.

The evidentiary record contradicts the statement on the basis of the official legislative record. The student loan forgiveness program was implemented through executive action under the HEROES Act, not through congressional legislation. No bill establishing the program passed either chamber of Congress. The legislative record is unambiguous and publicly documented.

The knowledge analysis establishes actual knowledge. The program was designed, developed, and announced by the administration through its own executive processes. The President cannot lack awareness of whether a program originating within his own administration was enacted through congressional legislation or executive action; the distinction is a basic feature of how the program came into existence and was understood within the administration from its inception.

Media Track

Example 5: Fox News and Dominion Voting Systems

Following the 2020 presidential election, Fox News broadcast repeated claims that Dominion Voting Systems machines had switched votes from Donald Trump to Joe Biden. These claims were made in broadcast capacity across multiple programs over an extended period. Each specific claim about machine vote-switching is a falsifiable factual assertion about the operation of voting technology, not a characterization or opinion about the election's integrity more broadly.

The evidentiary record contradicts the statements on the basis of official government data and an established record verified by multiple independent sources. The Department of Homeland Security's Cybersecurity and Infrastructure Security Agency issued a formal statement describing the 2020 election as the most secure in American history. State election officials, including Republicans in states where the claims were most prominently advanced, documented that the specific machine manipulation claims were unsupported by any evidence in the auditable record. No court reviewing challenges based on these claims found evidentiary support for them.

The knowledge analysis establishes actual institutional knowledge. Internal communications produced during the Dominion Voting Systems defamation litigation documented that Fox News hosts, producers, and senior executives possessed specific information contradicting the vote-switching claims during the period those claims were being broadcast. The documentary record established that individuals with editorial authority over the relevant broadcasts were aware that the specific factual claims being made were unsupported, and in several instances expressed that awareness explicitly in internal communications while the claims continued to air.

Example 6: CBS News and the National Guard documents

CBS News broadcast a report in September 2004 asserting that documents showed President George W. Bush had received preferential treatment in the Texas Air National Guard and had disobeyed a direct order. The central factual claim of the report was the authenticity of the documents on which it relied. Document authenticity is a specific, falsifiable factual claim subject to forensic analysis.

The evidentiary record contradicts the authenticity claim on the basis of an established record verified by multiple independent sources. Within days of broadcast, typography and forensic document experts identified specific characteristics of the documents — including font spacing and superscript formatting — inconsistent with typewriter technology available at the time the documents were purportedly produced, and consistent with modern word processing software. These findings were produced independently by multiple experts and were not subject to serious methodological dispute.

The knowledge analysis establishes constructive institutional knowledge. The forensic vulnerabilities in the documents were identifiable through analysis available to experts in the relevant field at the time of broadcast. CBS's own pre-broadcast verification process, as subsequently documented through the independent review commissioned by CBS itself, did not include the forensic document analysis that the outlet's editorial standards for authenticating historical documents would reasonably require. The post-broadcast period, during which CBS initially resisted the forensic findings before commissioning an independent review and ultimately retracting, documents the transition from constructive to actual institutional knowledge and the editorial response to that transition.

VII. Non-Qualifying Examples

The examples in this section illustrate statements that do not meet the methodology's standards for a finding, organized by the exclusion category that applies. They are as important to the methodology as the qualifying examples in Section VI, because the framework's credibility depends equally on what it does not reach as on what it does. Each example is drawn from real political discourse and is selected because it represents a category of statement the oversight body will encounter frequently — cases that may appear to qualify on first reading but fail the double filter or fall within an explicit exclusion.

The structure of each example identifies the statement, the exclusion category that applies, and the analytical reasoning that supports the exclusion. Where a statement fails on multiple grounds, the primary exclusion is identified and the secondary grounds are noted.

Forward-Looking Predictions

Example 1: "This tax cut will pay for itself"

Republican administrations and legislators have asserted on multiple occasions in official capacity settings that proposed tax reductions would generate sufficient economic growth to offset lost revenue. The statement is forward-looking: it asserts a future economic outcome contingent on a chain of causal mechanisms whose realization cannot be confirmed or refuted by evidence existing at the time of the statement.

This statement does not qualify for review under the forward-looking predictions exclusion. The existence of contrary Congressional Budget Office projections at the time of the statement does not transform the prediction into a falsifiable factual claim; expert forecasting, including official government forecasting, addresses future probabilities rather than present facts. The subsequent failure of the prediction to materialize does not retroactively satisfy the unambiguous evidentiary record condition, because the methodology applies the evidentiary standard at the time of the statement rather than in hindsight. The statement is a contestable economic prediction protected from review regardless of how confidently it was asserted or how consistently it has failed to prove accurate.

Example 2: Inflation described as "transitory"

Members of the Biden administration, including the President, characterized inflation in 2021 as transitory, implying it would resolve without sustained intervention. The characterization is forward-looking: it asserts a prediction about the trajectory of an economic phenomenon rather than a present fact about its current state.

This statement does not qualify for review under the forward-looking predictions exclusion for the same structural reasons as Example 1. The underlying disagreement among economists about inflation dynamics in 2021 further supports the exclusion under the contested scientific and empirical territory category, providing a secondary ground for non-qualification. The subsequent persistence of inflation does not establish that the characterization was false at the time it was made; it establishes only that the prediction did not materialize, which is categorically different from a finding of falsity under this methodology.

Opinion and Value Judgment

Example 3: "The economy is the greatest in history"

President Trump and other officials used variations of this formulation repeatedly in official capacity settings. The statement appears on first reading to assert a falsifiable comparative fact, but the comparative superlative depends entirely on which economic metrics are selected, how they are weighted, and across what time period the comparison is drawn. Reasonable analysts applying different but legitimate frameworks will reach different conclusions about which period in American economic history ranks highest by which measures.

This statement does not qualify for review under the opinion and value judgment exclusion. It fails the falsifiability condition because it does not have a determinable truth value independent of the assessor's choice of metrics and framework. It also fails under the comparative superlatives threshold category established in Section II, which identifies this class of statement as one that frequently resembles a falsifiable factual claim without meeting the standard. The statement is an evaluative characterization of economic conditions rather than a specific, verifiable factual assertion.

Example 4: "Republicans want to destroy Medicare"

Democratic politicians and officials have made variations of this claim in official capacity settings during debates over Medicare policy and federal budget negotiations. The statement attributes a specific intent or motivation to a class of political actors rather than asserting an observable fact about their conduct or the effects of their policies.

This statement does not qualify for review under the opinion and value judgment exclusion and independently fails under the statements of intent and attributed motivation exclusion established in Section III. Intent is not directly observable and cannot be confirmed or refuted by the evidentiary categories the methodology recognizes. The statement's apparent factual form — asserting what a political party "wants" — does not make it falsifiable; it remains a characterization of motivation regardless of how it is framed. Even substantial circumstantial evidence about the likely effects of proposed policies would not satisfy the unambiguous evidentiary record condition for a claim about intent.

Contested Scientific and Empirical Territory

Example 5: Attributing a specific weather event to climate change

Politicians and officials across both parties have at various times asserted in official capacity settings that specific hurricanes, wildfires, or other discrete weather events were caused by climate change. The broader scientific consensus on anthropogenic climate change is not in dispute and does not provide grounds for exclusion. The specific causal attribution of individual weather events to long-term climate trends, however, involves probabilistic reasoning in a field — climate attribution science — where the expert community has not resolved the relevant questions to the level of certainty the unambiguous threshold requires.

This statement does not qualify for review under the contested scientific and empirical territory exclusion. The exclusion does not apply to the underlying scientific consensus on climate change, which is sufficiently settled to support findings in appropriate cases. It applies specifically to the deterministic attribution of individual events, where the science remains probabilistic and the evidentiary record does not meet the unambiguous threshold. This distinction is important: the exclusion is narrow and does not provide cover for false statements about climate science generally.

Rhetorical Hyperbole

Example 6: "The most corrupt administration in history"

Politicians across both parties have used variations of this formulation in official capacity settings to characterize opposing administrations. The statement deploys a superlative comparative judgment in a context where its communicative function is emphatic and evaluative rather than literally assertive. A reasonable audience understands the statement as rhetorical characterization rather than a specific falsifiable claim about a measurable quantity.

This statement does not qualify for review under the rhetorical hyperbole exclusion. It fails the falsifiability condition on two independent grounds: it is a comparative superlative whose ranking depends on the assessor's choice of metrics, and it is rhetorical in register rather than literally assertive. The contextual standard established in Section III applies: the statement would not be understood by a reasonable audience as a literal factual assertion about a determinable historical ranking.

Scope Limitations

Example 7: Print and digital media — jurisdictional exclusion

In 2014, Rolling Stone magazine published an account of an alleged sexual assault at a University of Virginia fraternity that subsequently collapsed under journalistic scrutiny. A review conducted by the Columbia University Graduate School of Journalism documented systematic failures in the outlet's verification process. Rolling Stone retracted the story and settled litigation arising from it.

The specific factual claims in the original publication would warrant analysis under the double filter if made by a covered entity. Rolling Stone is not a covered entity. It is neither a public broadcaster receiving federal funding nor a licensed broadcaster operating on the public electromagnetic spectrum. It falls within the disclosure requirements tier of the media track, which does not include a findings mechanism. The oversight body has no jurisdiction over the publication's content regardless of the severity of the evidentiary failures involved.

This example is included not to minimize the significance of the publication's failures but to illustrate the jurisdictional boundary precisely. The same failures, had they occurred at a covered broadcaster, would have supported analysis under

Sections IV-B and V-B. The jurisdictional exclusion is not an evidentiary determination; it is a threshold determination that precedes evidentiary analysis entirely.

VIII. The Appeals Process

The appeals process is the primary institutional check on the oversight body's authority. Its function is not merely to protect the rights of officials and entities subject to findings; it is to create a structural incentive for the body to maintain evidentiary discipline by making every finding subject to specific, public challenge. A finding that does not survive a well-reasoned appeal is a visible indictment of the body's own rigor, and the architecture of the appeals process is designed to make that consequence real.

Standing

Primary standing to appeal a finding belongs to the official or broadcast entity that is its subject. Standing is automatic for the primary subject; no threshold showing of harm or interest is required beyond the existence of the finding itself.

Secondary standing is available to any party with a direct and documented interest in the outcome of the finding. This includes congressional ethics committees that have initiated or are considering proceedings on the basis of the finding, licensing authorities whose renewal determinations may be informed by it, and any official or entity whose own documented record is directly implicated by the finding's evidentiary reasoning. Secondary standing requires a threshold showing of direct interest filed within the appeal window; it is not available to parties whose interest is general or political rather than specific and procedural.

Standing is not available to parties seeking to challenge a finding on the grounds that the body should have issued a finding it did not issue. The appeals process governs adverse findings, not the body's exercise of discretion in declining to act. A party who believes the body failed to issue a warranted finding has recourse through the body's public transparency requirements and through existing political and judicial mechanisms, but not through the appeals process established here.

Grounds

Appeals are confined to three specific grounds. The specificity of these grounds is itself a structural feature: an open-ended appeals standard would allow the process to become a vehicle for relitigating the underlying political controversy rather than challenging the body's evidentiary reasoning.

The first ground is factual error. An appeal on this ground must identify a specific piece of evidence the body relied upon that is demonstrably incorrect, or a specific piece of contradicting evidence the body failed to consider that would materially affect the finding. A factual error appeal is not an opportunity to reargue the weight of the evidence; it requires identification of a specific error in the evidentiary record rather than a general disagreement with the body's conclusions.

The second ground is methodological deviation. An appeal on this ground must demonstrate that the body's analysis departed from the standards established in this methodology document in a specific and material way. This includes applying an evidentiary category not recognized by the methodology, failing to apply the correct knowledge standard for the relevant track, or reaching a finding on a claim that falls within an exclusion category established in Section III. A methodological deviation appeal requires identification of the specific provision of this document from which the body's analysis departed and an explanation of how that departure affected the outcome.

The third ground is procedural violation. An appeal on this ground must demonstrate that the body failed to follow the procedural requirements established in this document or in the oversight body's governing statute in a way that materially affected the finding. This includes failure to provide adequate notice, failure to allow a response period, or failure to publish the evidentiary basis of the finding in the form the methodology requires. A procedural violation appeal does not require a showing that the underlying finding is substantively incorrect; a material procedural failure is independently sufficient to vacate a finding regardless of the strength of the evidentiary record.

Timeline

Appeals must be filed within sixty days of the publication of a finding. The sixty-day window applies to both primary and secondary appellants; there is no

tolling of the window for secondary appellants pending the primary appellant's decision whether to appeal. Late filings are not accepted except upon a showing of extraordinary circumstance that prevented timely filing, which is assessed by the appellate body on a case by case basis.

Upon receipt of a timely filed appeal, the appellate body has thirty days to determine whether the appeal presents a claim with sufficient basis to warrant full review on at least one of the three recognized grounds. Appeals that do not meet this threshold are dismissed with a published explanation. Appeals that meet this threshold proceed to full review.

Full review is completed within one hundred and eighty days of the colorability determination. The appellate body may request additional submissions from the appellant and from the oversight body during this period. The oversight body is required to provide its full evidentiary record to the appellate body upon request; no portion of the record underlying a finding is privileged against appellate review.

The Appellate Body

The appellate body is structurally independent from the oversight body and is constituted under the same appointment principles: staggered long terms, supermajority confirmation requirements, and explicit exclusion of active or recent politicians, party officials, and media executives. Its members are drawn from the legal and academic communities, with particular emphasis on expertise in evidentiary standards, administrative law, and the relevant substantive domains.

The appellate body does not conduct de novo review of the underlying factual record. Its function is to assess whether the oversight body's analysis met the standards this methodology establishes, not to substitute its own judgment about the weight of the evidence. On factual error grounds, the appellate body determines whether the identified error is real and material. On methodological deviation grounds, it determines whether the body departed from the established standards and whether that departure affected the outcome. On procedural violation grounds, it determines whether the violation occurred and whether it was material.

The appellate body has three available dispositions. It may affirm the finding, in which case the finding stands and the affirmation is published alongside it. It

may vacate the finding, in which case the finding is removed from the public record and replaced with a published explanation of the grounds for vacation. It may remand the finding to the oversight body for reconsideration consistent with the appellate body's published reasoning, in which case the oversight body has ninety days to issue a revised finding or withdraw the original.

Publication Requirements

Every stage of the appeals process is published in full and in real time. This includes the initial appeal filing and its stated grounds, the colorability determination and its reasoning, all submissions made during full review, the appellate body's final disposition and its complete reasoning, and in remand cases, the oversight body's response and any revised finding.

The publication requirement serves two functions simultaneously. It subjects the appellate body's own reasoning to public scrutiny, preventing the appeals process from becoming a second venue for opaque institutional decision-making. And it ensures that a vacated finding is visible as a vacated finding rather than simply disappearing from the record; the public record reflects not only what the oversight body found but what the appellate body determined about the quality of that finding.

A finding that has been vacated on appeal is removed from the active public record but is not expunged. The full record of the original finding, the appeal, and the vacation is preserved in an archival record that is publicly accessible. This ensures that the oversight body's errors are part of the permanent institutional record rather than hidden, creating a longitudinal accountability mechanism for the body's own performance over time.

IX. Hard Cases

The methodology's value is tested not by its application to clear cases but by its handling of genuinely difficult ones. This section addresses the category of claims that resist straightforward analysis — cases where the evidentiary record is substantial but contested, where the knowledge standard is difficult to apply, or where the double filter's boundaries require careful reasoning rather than mechanical application. The purpose is not to resolve these cases definitively but

to demonstrate how the methodology's analytical framework applies when the facts are genuinely complex.

Hard cases do not represent failures of the methodology; they represent its most important function. A methodology that produces clean outcomes in every case is either operating only on trivial claims or papering over genuine analytical difficulty. The framework's credibility depends on its capacity to reason transparently through difficult cases, including cases where the conclusion is that no finding is warranted despite a substantial public record of apparent falsity.

The Iraq WMD Claims

The claims made by Bush administration officials in the period preceding the 2003 invasion of Iraq — including Secretary of State Colin Powell's presentation to the United Nations Security Council — represent the most analytically complex case the politician track's knowledge standard is likely to encounter in recent American history. The case is included here not because it produces a clean finding but because working through it honestly demonstrates the methodology's limits as clearly as its application.

The statements at issue were specific and falsifiable: that Iraq possessed active weapons of mass destruction programs, that specific facilities were producing prohibited materials, and that the evidence for these claims was solid. These were not forward-looking predictions or evaluative characterizations; they were present-tense factual assertions made in official capacity settings of the highest profile.

The evidentiary record is where the analysis becomes genuinely difficult. The subsequent failure to find the asserted weapons programs is not itself sufficient to establish that the statements were false at the time they were made under the methodology's standards. The relevant question is not what was found after the invasion but what the evidentiary record established at the time of the statements. The intelligence assessments available to administration officials at the time reflected a genuine, if subsequently discredited, analytical consensus within the intelligence community. The October 2002 National Intelligence Estimate, the document most directly relevant to the knowledge analysis, reflected that consensus while also containing dissenting assessments from certain agencies that were not prominently reflected in public statements.

This is where the hard case analysis turns. The knowledge standard requires asking whether administration officials possessed information contradicting their public assertions, or whether the contradicting information was sufficiently available that they reasonably should have known. The existence of dissenting intelligence assessments within the overall NIE is relevant but not clearly sufficient to establish constructive knowledge under the methodology's standard, because the dissents were minority positions within the official intelligence community's own documented consensus at the time. A finding would require establishing that specific officials had access to specific contradicting assessments and made public statements inconsistent with those assessments — a showing the public record does not cleanly support for most of the relevant statements, though it may support specific claims made by specific officials in specific settings.

The harder analytical problem is that the intelligence community's own consensus proved wrong. The methodology does not provide a mechanism for finding falsity based on an evidentiary record that was itself incorrect at the time; the four authoritative evidentiary categories require that the contradicting evidence be reliable, and intelligence assessments that were themselves fabricated or analytically flawed do not satisfy that standard retroactively. The framework reaches deliberate or negligent misrepresentation, not institutional analytical failure of the kind the Iraq intelligence record ultimately represents.

The conclusion the methodology supports is neither a clean finding nor a clean non-finding. Some specific claims made by specific officials in specific settings may support finding analysis on constructive knowledge grounds, particularly where public statements went beyond what the underlying intelligence assessments actually established. The general category of WMD assertions, however, does not support a finding under the knowledge standard as applied to most officials making most statements during the relevant period. The appropriate disposition for this category of case is a documented record of the analytical difficulty rather than a forced finding in either direction.

This conclusion will be unsatisfying to those who believe the administration knowingly misled the public about the case for war. The methodology does not adjudicate that question; it applies a specific evidentiary standard to specific claims, and the evidentiary record as publicly documented does not clearly satisfy that standard for most of the relevant statements. This is precisely the kind of

outcome that demonstrates the methodology is functioning as designed rather than as a vehicle for reaching predetermined conclusions.

The Boundary Between Prediction and Present Fact

A second category of hard case arises when a statement straddles the boundary between a forward-looking prediction and a present-tense factual assertion. The exclusion for forward-looking predictions is clear in its core application but requires careful reasoning at its edges.

Consider a statement by an official that a specific policy "has already created" a specific number of jobs. The present perfect tense appears to assert a present fact rather than a future prediction, making it potentially falsifiable against employment data existing at the time of the statement. But job creation figures are themselves subject to revision, methodological dispute among economists, and genuine uncertainty about causation — which metrics capture which jobs, and which of those jobs are attributable to which policy. The statement has the grammatical form of a falsifiable factual claim but the underlying evidentiary situation of a contested empirical interpretation.

The double filter's application to this category requires asking two questions in sequence. First, is the specific numerical claim falsifiable against an unambiguous data source — for instance, does the Bureau of Labor Statistics publish a figure that directly contradicts the number asserted? If yes, the falsifiability condition may be met. Second, is the causal attribution — that the policy created those jobs — itself falsifiable to the unambiguous threshold, or does it involve the kind of contested economic methodology that falls within the contested scientific and empirical territory exclusion? The answer to the second question will typically be no, because causal attribution in macroeconomics involves genuine methodological dispute even when the underlying employment figures are not in dispute.

The practical result is that specific numerical claims about present economic conditions may qualify for review when they are directly contradicted by official government data, but causal attribution claims bundled with those numbers typically will not. The finding, if any, attaches to the specific numerical assertion and not to the causal attribution, consistent with the principle established in Section III that the body may review the factual component of a bundled claim while leaving the evaluative or contested component outside its jurisdiction.

The Evolving Statement Problem

A third category of hard cases arises when an official's position on a factual matter evolves over time in ways that create apparent contradictions between earlier and later statements. The self-contradiction category in Section V-A establishes that prior documented statements in official capacity are an evidentiary input, and that a statement contradicting a prior official statement raises a presumption of actual knowledge. But the presumption is rebuttable where the official's position changed on the basis of new information.

The hard case is distinguishing a genuine evidentiary update from a post-hoc rationalization of a position change driven by political rather than evidentiary considerations. The methodology does not provide a mechanical answer to this question, but it does provide an analytical framework. A genuine evidentiary update should be identifiable in the record: a specific new study, a revised government data release, or a documented change in expert consensus that postdates the prior statement and is consistent with the direction of the position change. A position change that cannot be anchored to a specific evidentiary development, or that occurs in a context where political incentives clearly point in the same direction as the change, warrants closer scrutiny under the knowledge standard without automatically satisfying it.

The oversight body is not equipped to adjudicate political motivation, and the methodology does not ask it to. What it can do is document the evidentiary record surrounding a position change and assess whether the record supports the inference of a genuine update. Where it does not, the prior statement remains available as an evidentiary input and the knowledge analysis proceeds accordingly.

Conclusion

This document and the white paper it accompanies are designed to be read as a single integrated framework. The white paper establishes the legal architecture and constitutional theory; this document operationalizes that architecture by defining with precision the standards that govern every finding the oversight body produces. Neither document is complete without the other, and the

constitutional defensibility of the framework depends in material part on the rigor of the methodology described here.

The methodology's central purpose is to draw a principled, transparent, and consistently defensible line between the vast domain of protected political speech and the narrow category of empirically falsifiable claims supported by an unambiguous evidentiary record. That line is not drawn by institutional discretion or political judgment; it is drawn by the double filter, the exclusion categories, the evidentiary standards, and the knowledge requirements established in the preceding sections. Every element of the methodology serves this single purpose: ensuring that the oversight body's findings are the product of a reproducible process rather than a reflection of the body's own preferences.

The framework accepts that this line will sometimes produce outcomes that feel incomplete or unsatisfying. The Iraq WMD analysis in Section IX reaches a conclusion that many will find inadequate given the scale of the consequences that followed from those statements. The forward-looking predictions exclusion protects economic claims that proved spectacularly wrong. The contested scientific territory exclusion declines to adjudicate disputes that feel resolved to many observers. These outcomes are not failures of the methodology; they are evidence that it is functioning as designed. A methodology that reaches every case where public opinion believes falsity occurred is not a neutral evidentiary standard — it is a mechanism for institutionalizing the conclusions of whoever controls it. The discipline to decline jurisdiction where the evidentiary standard is not met is the methodology's most important feature, not its most significant limitation.

The framework also accepts that it will be attacked from both directions simultaneously. Those who believe it reaches too far will argue that any institutional mechanism for documenting official false statements is a threat to free expression regardless of how narrowly it is drawn. Those who believe it does not reach far enough will argue that the evidentiary standards are too demanding, the exclusions too broad, and the consequences too modest to matter. Both critiques are predictable, and neither is answered by adjusting the methodology in response. The answer is the methodology itself: its transparency, its specificity, its symmetrical application across partisan lines, and its willingness to reach conclusions that are uncomfortable for both sides.

The methodology is a living document in one specific sense: the worked examples will be updated as the oversight body's findings accumulate and as the appellate

record develops. The analytical framework established here — the double filter, the exclusion categories, the evidentiary standards, the knowledge requirements, and the appeals architecture — is not subject to revision through the examples update process. Changes to the analytical framework require the same public process and cross-partisan consensus that governs the oversight body's own operations. This distinction between the stable analytical core and the evolving evidentiary record is itself a structural safeguard: it ensures the methodology cannot be gradually revised through the accumulation of precedent in ways that expand the body's jurisdiction beyond the boundaries established here.

The work of building institutional credibility begins before any formal authority is sought. The accountability tracker, the advisory council, and the public documentation of the methodology's application in practice are the proof of concept on which that credibility depends. The question the framework ultimately asks is not whether official falsehood is a problem worth addressing — that question has already been answered by the scale of the damage it has caused. The question is whether American democratic institutions are capable of building the oversight architecture that addresses it with the rigor, the discipline, and the symmetry the problem demands. This methodology is the answer to that question in the only form that matters: not a promise, but a procedure.